

AI-BASED DIABETES PREDICTION SYSTEM USING MACHINE LEARNING

TEAM MEMBER

922321106008: S.DEVIGA

PROJECT TITLE: AI-based diabetes prediction system

PHASE4: Development part 2

TOPIC: AI-based diabetes prediction model



INTRODUCTION:

Creating an AI-based diabetes prediction system involves several key steps: feature engineering, model training, and evaluation. Here's a brief introduction to each of these components:

FEATURE ENGINEERING:

Feature engineering is the process of selecting, transforming, and creating relevant features (input variables) from the available data that will be used to train your AI model. In the context of a diabetes prediction system, some common features might include age, BMI, family medical history, glucose levels, and more. Feature engineering also involves handling missing data, scaling, and encoding categorical variables, making the data suitable for machine learning.

MODEL TRAINING:

Model training involves the development of a machine learning or deep learning model using your dataset. In the case of diabetes prediction, you might use algorithms like logistic regression, decision trees, random forests, support vector machines, or neural networks. The model learns patterns from the features in your dataset and makes predictions.

EVALUATION:

After training your model, you need to assess its performance. Common evaluation metrics for a binary classification problem like diabetes prediction include accuracy, precision, recall, F1 score, and area under the ROC curve (AUC-ROC). You'll typically split your dataset into a training set and a test set to evaluate how well your model generalizes to new, unseen data. Cross-validation can also be used to assess the model's robustness.

It's important to note that the choice of features, the model architecture, and the evaluation metrics may vary depending on the specific requirements and characteristics of your dataset and application. Continuous refinement and fine-tuning of these components are often necessary to build an accurate and reliable diabetes prediction system.

DATASETLINK:

<https://in.docworkspace.com/d/sIJjqyObiAZzV5KkG>

769 Rows*9 Columns

GIVEN DATASET:

pregnancies	Glucose	Blood Pressure	Skin Thickness	Insulin	BMI	Diabetes Pedigree Function	Age	Outcome
6	148	72	35	0	33.6	0.627	50	1
1	85	66	29	0	26.6	0.351	31	0
8	183	64	0	0	23.3	0.672	32	1
1	89	66	23	94	28.1	0.167	21	0
0	137	40	35	168	43.1	2.288	33	1
5	116	74	0	0	25.6	0.201	30	0
3	78	50	32	88	31	0.248	26	1
10	115	0	0	0	35.3	0.134	29	0
2	197	70	45	543	30.5	0.158	53	1
8	125	96	0	0	0	0.232	54	1
4	110	92	0	0	37.6	0.191	30	0
10	168	74	0	0	38	0.537	34	1
10	139	80	0	0	27.1	1.441	57	0
1	189	60	23	846	30.1	0.398	59	1
5	166	72	19	175	25.8	0.587	51	1
7	100	0	0	0	30	0.484	32	1
0	118	84	47	230	45.8	0.551	31	1

Here is the further procedure to proceed after data collection and data processing process.

Model selection: It is an AI-based diabetes prediction system that involves choosing the most appropriate machine learning or statistical model to achieve accurate predictions. Here's a simplified process:

Data Collection: Gather a dataset containing features (e.g., age, BMI, glucose levels) and target labels (diabetes status).

Data Preprocessing: Clean and preprocess the data by handling missing values, scaling features, and encoding categorical variables.

Split Data: Divide the dataset into training, validation, and test sets to evaluate model performance properly.

Model Selection: Consider various machine learning models such as logistic regression, decision trees, random forests, support vector machines, neural networks, and more. Factors to consider:

- Complexity of the model.
- Interpretability of the model.
- Performance metrics (accuracy, precision, recall, F1-score) on the validation set.
- Computational resources required.

Hyperparameter Tuning: For selected models, fine-tune hyperparameters (e.g., learning rate, depth of trees, number of hidden layers) using techniques like grid search or random search.

CROSS-VALIDATION: Perform k-fold cross-validation to ensure the model's generalization ability and reduce overfitting.

EVALUATE PERFORMANCE: Assess models on the test set using appropriate evaluation metrics for diabetes prediction. Consider the trade-offs between false positives and false negatives, as it depends on the application's requirements.

ENSEMBLE METHODS: Consider using ensemble techniques like bagging (e.g., Random Forests) or boosting (e.g., AdaBoost, Gradient Boosting) to combine multiple models for improved performance.

MODEL INTERPRETABILITY: Depending on the application, choose models that provide interpretability, such as decision trees or linear regression, to understand the factors influencing predictions.

DEPLOYMENT: Once you've selected the best-performing model, deploy it in your diabetes prediction system for real-world use.

PROGRAM:

INPUT:

```
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
scaler.fit(X)
SSX = scaler.transform(X)
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(SSX, y, test_size=0.2,
random_state=7)
```

```
X_train.shape, y_train.shape
```

OUTPUT:

```
((614, 8), (614,))
```

INPUT

```
X_test.shape, y_test.shape
```

OUTPUT

```
((154, 8), (154,))
```

Logistic Regression: Logistic regression can be a valuable tool in an AI-based diabetes prediction system

- Use logistic regression to build a predictive model.

Logistic regression is well-suited for binary classification task. Assess the performance of your logistic regression model using metrics like accuracy, precision, recall, F1-score, and ROC AUC. Cross-validation can help estimate model robustness.

- Consider using L1 or L2 regularization to prevent overfitting and improve model generalization.

- Logistic regression models are inherently interpretable, making it easier to explain to stakeholders and healthcare professionals why a certain prediction was made

CLASSIFICATION ALGORITHMS:

LOGISTIC REGRESSION:

INPUT

```
from sklearn.linear_model import LogisticRegression
lr = LogisticRegression(solver='liblinear', multi_class='ovr')
lr.fit(X_train, y_train)
```

OUTPUT

```
LogisticRegression
LogisticRegression(multi_class='ovr', solver='liblinear')
```

Decision Tree: use a decision tree algorithm to build a predictive model. Decision trees are useful for both classification and regression tasks. In this case, you would use it for binary classification (diabetes or no diabetes).

- Decision trees can automatically rank feature importance, helping you understand which factors contribute the most to diabetes prediction.
- Assess the performance of your decision tree model using metrics like accuracy, precision, recall, F1-score, and ROC AUC. Cross-validation can help estimate model robustness.

- Decision trees can be prone to overfitting, so consider pruning the tree to improve generalization and avoid making overly complex decisions. Combine multiple decision trees using ensemble methods like Random Forests or Gradient Boosting to enhance prediction accuracy and reduce overfitting.
- Decision trees are inherently interpretable, making it easy to explain the reasoning behind a prediction to healthcare professionals and patients.

MAKING PREDICTION:

Logistic Regression:

INPUT

```
X_test.shape
```

OUTPUT

```
(154, 8)
```

INPUT

```
lr_pred=lr.predict(X_test)
```

```
lr_pred.shape
```

OUTPUT

```
(154,)
```

Decision Tree:

```
dt_pred=dt.predict(X_test)
```

```
dt_pred.shape
```

OUTPUT

```
(154,)
```

Accuracy: Accuracy is a fundamental performance metric in an AI-based diabetes prediction system. It provides an overall

measure of how well the model is performing. Here's how accuracy is relevant to such a system:

Accuracy is calculated as:

$$\text{Accuracy} = \frac{\text{correct predictions}}{\text{total predictions}}$$

Correct Predictions: The sum of true positives (correctly predicted cases of diabetes) and true negatives (correctly predicted cases of non-diabetes).

Total Predictions: The total number of predictions made by the model.

➤ A high accuracy indicates that the model is making a large proportion of correct predictions, both for patients with diabetes and those without it. It's a straightforward and easy-to-understand metric that can provide an initial assessment of the model's performance.

However, it's important to note that accuracy might not be the only relevant metric, especially in healthcare applications. Diabetes prediction systems often deal with imbalanced datasets, where non-diabetic cases significantly outnumber diabetic cases. In such cases, a model can achieve high accuracy by simply predicting "no diabetes" for every case, but this would not be clinically useful.

To get a more comprehensive evaluation, consider using additional metrics like precision, recall, and the F1-score, which provide insights into the model's ability to correctly

predict diabetes cases and non-diabetes cases while accounting for imbalances in the dataset

MODEL EVALUATION FOR LOGISTIC REGRESSION:

INPUT

```
from sklearn.metrics import accuracy_score
print("Train Accuracy of Logistic Regression: ", lr.score(X_train,
y_train)*100)
print("Accuracy (Test) Score of Logistic Regression: ", lr.score(X_test,
y_test)*100)
print("Accuracy Score of Logistic Regression: ", accuracy_score(y_test,
lr_pred)*100)
```

OUTPUT

```
Train Accuracy of Logistic Regression: 77.36156351791531
Accuracy (Test) Score of Logistic Regression: 77.27272727272727
Accuracy Score of Logistic Regression: 77.27272727272727
```

MODEL EVALUATION FOR DECISION TREE

INPUT

```
print ("Train Accuracy of Decesion Tree: ", dt.score(X_train, y_train)*100)
print("Accuracy (Test) Score of Decesion Tree: ", dt.score(X_test, y_test)*100)
print("Accuracy Score of Decesion Tree: ", accuracy_score(y_test, dt_pred)*100)
```

OUTPUT

```
Train Accuracy of Decesion Tree: 100.0
Accuracy (Test) Score of Decesion Tree: 81.16883116883116
Accuracy Score of Decesion Tree: 81.16883116883116
```

Confusion matrix: A confusion matrix is a valuable tool in assessing the performance of an AI-based diabetes prediction system. It provides a detailed breakdown of the model's predictions, allowing you to understand how well it's classifying diabetes and non-diabetes cases. The confusion

matrix is especially useful in the context of binary classification, where you have two classes: diabetes (positive) and non-diabetes (negative).

A confusion matrix consists of four key components:

True Positives (TP): The number of cases correctly predicted as having diabetes .

True Negatives (TN): The number of cases correctly predicted as not having diabetes.

False Positives (FP): The number of cases incorrectly predicted as having diabetes when they do not (Type I error).

False Negatives (FN): The number of cases incorrectly predicted as not having diabetes when they do have diabetes (Type II error).

With this information, you can calculate various performance metrics, such as:

Accuracy: $TP+TN/TP+TN+FP+FN$

Precision: $TP/TP+FP$

Recall(sensitivity): $TP/TP+FN$

Specificity: $TN/TN+FP$

F1-Score: A measure that balances precision and recall, computed as

$2 \cdot \text{precision} \cdot \text{recall} / (\text{precision} + \text{recall})$

The confusion matrix and these associated metrics provide a comprehensive view of your model's performance, helping you understand its ability to correctly predict diabetes cases while minimizing false alarms or missed cases. This information is crucial for assessing the clinical utility of your diabetes prediction system.

INPUT

```
from sklearn.metrics import classification_report, confusion_matrix
cm = confusion_matrix(y_test, lr_pred)
cm
```

OUTPUT


```
array([[86, 11],
       [24, 33]])
```

INPUT

```
sns.heatmap(confusion_matrix(y_test, lr_pred), annot=True, fmt="d")
```

OUTPUT

<Axes: >



INPUT

```
TN =cm[0, 0]
FP =cm[0,1]
FN = cm[1,0]
TP = cm[1,1]
```

TN, FP, FN, TP

OUTPUT

(86, 11, 24, 33)

INPUT

```
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.metrics import accuracy_score, roc_auc_score, roc_curve
cm = confusion_matrix(y_test, lr_pred)

print('TN - True Negative {}'.format(cm[0,0]))
print('FP - False Positive {}'.format(cm[0,1]))
print('FN - False Negative {}'.format(cm[1,0]))
```



```

print('TP - True Positive {}'.format(cm[1,1]))
print('Accuracy Rate: {}'.format(np.divide(np.sum([cm[0,0], cm[1,1]]),
np.sum(cm))*100))
print('Misclassification Rate: {}'.format(np.divide(np.sum([cm[0,1], cm[1,0]]),
np.sum(cm))*100))

```

OUTPUT

```

TN - True Negative 86
FP - False Positive 11
FN - False Negative 24
TP - True Positive 33
Accuracy Rate: 77.27272727272727
Misclassification Rate: 22.727272727272727

```

INPUT

```
77.27272727272727+22.727272727272727
```

OUTPUT

```
100.0
```

INPUT

```


import matplotlib.pyplot as plt
import numpy as np

plt.clf()
plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Wistia)
classNames = ['0', '1']
plt.title('Confusion Matrix of Logistic Regression')
plt.ylabel('Actual (true) Values')
plt.xlabel('Predicted Values')
tick_marks = np.arange(len(classNames))
plt.xticks(tick_marks, classNames, rotation=45)
plt.yticks(tick_marks, classNames)
s = [['TN', 'FP'], ['FN', 'TP']]
for i in range(2):
    for j in range(2):
        plt.text(j, i, str(s[i][j]) + " = " + str(cm[i][j]))

plt.show()

```

OUTPUT



INPUT

```
pd.crosstab(y_test, lr_pred, margins=False)
```

OUTPUT

col_0	0	1
Outcome		
0	86	11
1	24	33

INPUT

```
pd.crosstab(y_test, lr_pred, margins=True)
```

OUTPUT

col_0	0	1	All
Outcome			
0	86	11	97
1	24	33	57
All	110	44	154

INPUT

```
pd.crosstab(y_test, lr_pred, rownames=['Actual values'], colnames=['Predicted values'], margins=True)
```

OUTPUT

Predicted values	0	1	All
Actual values			
0	86	11	97
1	24	33	57
All	110	44	154

Precision:

PPV- positive Predictive Value

Precision = True Positive / (True Positive + False Positive)

Precision = TP / (TP + FP)

INPUT

TP, FP

OUTPUT

(33, 11)

INPUT

Precision = TP/(TP+FP)
Precision

OUTPUT

0.75

INPUT

33 / (33+11)

OUTPUT

0.75

precision Score:

INPUT

```
precision_score = TP/float(TP+FP)*100  
print('Precision Score: {0:0.4f}'.format(precision_score))
```

OUTPUT

Precision Score: 75.0000

INPUT

```
from sklearn.metrics import precision_score  
print("Precision Score is: ", precision_score(y_test, lr_pred)*100)  
print("Micro Average Precision Score is: ", precision_score(y_test, lr_pred,  
average='micro')*100)  
print("Macro Average Precision Score is: ", precision_score(y_test, lr_pred,  
average='macro')*100)  
print("Weighted Average Precision Score is: ", precision_score(y_test, lr_pred,  
average='weighted')*100)  
print("precision Score on Non Weighted score is: ", precision_score(y_test,  
lr_pred, average=None)*100)
```

OUTPUT

Precision Score is: 75.0

Micro Average Precision Score is: 77.27272727272727

Macro Average Precision Score is: 76.590909090909091

Weighted Average Precision Score is: 77.00413223140497

precision Score on Non Weighted score is: [78.18181818 75.]

INPUT

```
print('Classification Report of Logistic Regression: \n',  
classification_report(y_test, lr_pred, digits=4))
```

OUTPUT

Classification Report of Logistic Regression:

	precision	recall	f1-score	support
0	0.7818	0.8866	0.8309	97
1	0.7500	0.5789	0.6535	57
accuracy			0.7727	154
macro avg	0.7659	0.7328	0.7422	154
weighted avg	0.7700	0.7727	0.7652	154

Precision:

Precision is an important performance metric in an AI-based diabetes prediction system. It measures the accuracy of positive predictions made by the model. Here's how precision is relevant to such a system:

Precision is the ratio of true positive predictions to the total positive predictions made by the model. In the context of a diabetes prediction system:

True Positives (TP): These are cases where the model correctly predicts a patient as having diabetes, and they indeed have diabetes.

False Positives (FP): These are cases where the model incorrectly predicts a patient as having diabetes, but they do not have diabetes.

Precision is calculated as:

$\text{True positives} / (\text{true positives} + \text{false positives})$

➤ A high precision indicates that when the model predicts a patient has diabetes, it's more likely to be correct. In a healthcare setting, high precision is crucial because it means fewer false alarms or unnecessary treatments. It helps in avoiding unnecessary stress and expenses for patients and healthcare resources.

However, precision should be considered in conjunction with other metrics like recall and F1-score, as there's often a trade-off between precision and recall. Achieving a balance between precision and recall is essential to ensure that the model correctly identifies diabetes cases while minimizing false alarms.

True Positive Rate(TPR)

$\text{Recall} = \text{True Positive} / (\text{True Positive} + \text{False Negative})$

$\text{Recall} = \text{TP} / (\text{TP} + \text{FN})$

INPUT

```
recall_score = TP / float(TP+FN)*100  
print('recall_score', recall_score)
```

OUTPUT

recall_score 57.89473684210527

INPUT

TP, FN

OUTPUT

(33, 24)

INPUT

$33 / (33 + 24)$

OUTPUT

0.5789473684210527

INPUT

```
from sklearn.metrics import recall_score
print('Recall or Sensitivity_Score: ', recall_score(y_test, lr_pred)*100)
```

Recall or Sensitivity_Score: 57.89473684210527

```
print("recall Score is: ", recall_score(y_test, lr_pred)*100)
print("Micro Average recall Score is: ", recall_score(y_test, lr_pred,
average='micro')*100)
print("Macro Average recall Score is: ", recall_score(y_test, lr_pred,
average='macro')*100)
print("Weighted Average recall Score is: ", recall_score(y_test, lr_pred,
average='weighted')*100)
print("recall Score on Non Weighted score is: ", recall_score(y_test, lr_pred,
average=None)*100)
```

OUTPUT

recall Score is: 57.89473684210527

Micro Average recall Score is: 77.27272727272727

Macro Average recall Score is: 73.27726532826912

Weighted Average recall Score is: 77.27272727272727

recall Score on Non Weighted score is: [88.65979381 57.89473684]

INPUT

```
print('Classification Report of Logistic Regression: \n',
classification_report(y_test, lr_pred, digits=4))
```

OUTPUT

Classification Report of Logistic Regression:

	precision	recall	f1-score	support
0	0.7818	0.8866	0.8309	97
1	0.7500	0.5789	0.6535	57
accuracy			0.7727	154
macro avg	0.7659	0.7328	0.7422	154
weighted avg	0.7700	0.7727	0.7652	154

FPR - False Positive Rate

INPUT

```
FPR = FP / float(FP + TN) * 100
print('False Positive Rate: {:.4f}'.format(FPR))
```

OUTPUT

False Positive Rate: 11.3402

INPUT

FP, TN

OUTPUT

(11, 86)

INPUT

11/(11+86)

OUTPUT

0.1134020618556701

Recall: Recall, also known as sensitivity or true positive rate, is an important performance metric in an AI-based diabetes prediction system. It measures the model's ability to correctly identify all positive cases, or in this context, patients with diabetes. Here's how recall is relevant to such a system:

True Positives (TP): These are cases where the model correctly predicts a patient as having diabetes, and they indeed have diabetes.

False Negatives (FN): These are cases where the model incorrectly predicts a patient as not having diabetes, but they do have diabetes.

Recall is calculated as

Recall: $\text{True positives} / (\text{True positives} + \text{false positives})$

A high recall indicates that the model is effective at identifying patients with diabetes, minimizing the risk of missing true cases. In the context of a diabetes prediction system, high recall is crucial because it means the system is better at identifying individuals who need medical attention or diabetes management.

However, recall should be considered in conjunction with other metrics like precision and the F1-score. There's often a trade-off between precision and recall. While high recall minimizes false negatives, it may result in more false positives. Achieving a balance between precision and recall is

essential to ensure that the model correctly identifies diabetes cases while minimizing incorrect predictions.

INPUT

```
specificity = TN / (TN+FP)*100  
print('Specificity : {0:0.4f}'.format(specificity))
```

OUTPUT

Specificity : 88.6598

INPUT

```
from sklearn.metrics import f1_score  
print('F1_Score of Macro: ', f1_score(y_test, lr_pred)*100)
```

OUTPUT

F1_Score of Macro: 65.34653465346535

INPUT

```
print("Micro Average f1 Score is: ", f1_score(y_test, lr_pred,  
average='micro')*100)  
print("Macro Average f1 Score is: ", f1_score(y_test, lr_pred,  
average='macro')*100)  
print("Weighted Average f1 Score is: ", f1_score(y_test, lr_pred,  
average='weighted')*100)  
print("f1 Score on Non Weighted score is: ", f1_score(y_test, lr_pred,  
average=None)*100)
```

OUTPUT

Micro Average f1 Score is: 77.27272727272727
Macro Average f1 Score is: 74.21916104653944
Weighted Average f1 Score is: 76.52373933045479
f1 Score on Non Weighted score is: [83.09178744 65.34653465]

Classification Report of Logistic Regression:

```
from sklearn.metrics import classification_report
print('Classification Report of Logistic Regression: \n',
classification_report(y_test, lr_pred, digits=4))
```

Classification Report of Logistic Regression:

	precision	recall	f1-score	support
0	0.7818	0.8866	0.8309	97
1	0.7500	0.5789	0.6535	57
accuracy			0.7727	154
macro avg	0.7659	0.7328	0.7422	154
weighted avg	0.7700	0.7727	0.7652	154

ROC Curve& ROC AUC

INPUT

```
auc= roc_auc_score(y_test, lr_pred)
print("ROC AUC SCORE of logistic Regression is ", auc)
```

ROC AUC SCORE of logistic Regression is 0.7327726532826913

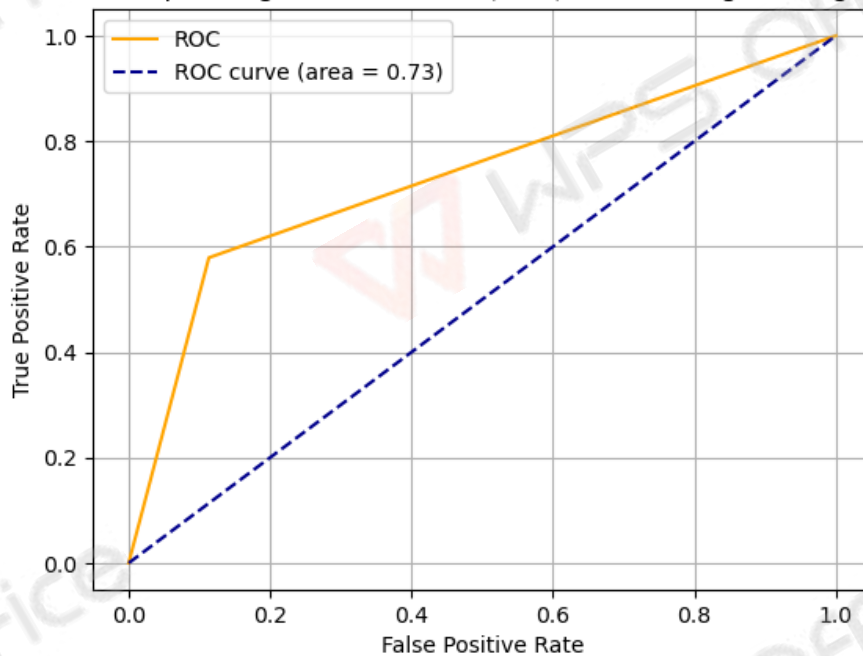
INPUT

```
from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt
```

```
fpr, tpr, thresholds = roc_curve(y_test, lr_pred)
plt.plot(fpr, tpr, color='orange', label="ROC")
plt.plot([0, 1], [0, 1], color='darkblue', linestyle='--', label='ROC curve
(area = %0.2f)' % auc(fpr, tpr))
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("Receiver Operating Characteristics (ROC) Curve of Logistic
Regression")
plt.legend()
plt.grid()
plt.show()
```

OUTPUT

Receiver Operating Characteristics (ROC) Curve of Logistic Regression



Confusion Matrix of Decision tree:

Confusion Matrix:

```
from sklearn.metrics import classification_report, confusion_matrix
```

```
cm = confusion_matrix(y_test, dt_pred)  
cm
```

OUTPUT


```
array([[83, 14],  
       [15, 42]])
```

INPUT

```
sns.heatmap(confusion_matrix(y_test, dt_pred), annot=True, fmt="d")
```

OUTPUT:

<Axes: >



INPUT

```
TN =cm[0, 0]  
FP =cm[0,1]  
FN = cm[1,0]  
TP = cm[1,1]
```

INPUT

TN, FP, FN, TP

OUTPUT

(83, 14, 15, 42)

INPUT

```
from sklearn.metrics import classification_report, confusion_matrix  
from sklearn.metrics import accuracy_score, roc_auc_score, roc_curve  
cm = confusion_matrix(y_test, dt_pred)
```

```

print('TN - True Negative {}'.format(cm[0,0]))
print('FP - False Positive {}'.format(cm[0,1]))
print('FN - False Negative {}'.format(cm[1,0]))
print('TP - True Positive {}'.format(cm[1,1]))
print('Accuracy Rate: {}'.format(np.divide(np.sum([cm[0,0], cm[1,1]]),
np.sum(cm))*100))
print('Misclassification Rate: {}'.format(np.divide(np.sum([cm[0,1],
cm[1,0]]), np.sum(cm))*100))

```

OUTPUT

```

TN - True Negative 83
FP - False Positive 14
FN - False Negative 15
TP - True Positive 42
Accuracy Rate: 81.16883116883116
Misclassification Rate: 18.83116883116883

```

INPUT

```


import matplotlib.pyplot as plt
import numpy as np

plt.clf()
plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Wistia)
classNames = ['0', '1']
plt.title('Confusion Matrix of Decision Tree')
plt.ylabel('Actual (true) Values')
plt.xlabel('Predicted Values')
tick_marks = np.arange(len(classNames))
plt.xticks(tick_marks, classNames, rotation=45)
plt.yticks(tick_marks, classNames)
s = [['TN', 'FP'], ['FN', 'TP']]
for i in range(2):
    for j in range(2):
        plt.text(j, i, str(s[i][j]) + " = " + str(cm[i][j]))

plt.show()

```

OUTPUT



```
specificity = TN / (TN+FP)*100  
print('Specificity : {0:0.4f}'.format(specificity))
```

Specificity : 88.6598

INPUT

```
from sklearn.metrics import f1_score  
print('F1_Score of Macro: ', f1_score(y_test, lr_pred)*100)
```

OUTPUT

F1_Score of Macro: 65.34653465346535

INPUT

```

print("Micro Average f1 Score is: ", f1_score(y_test, lr_pred,
average='micro')*100)
print("Macro Average f1 Score is: ", f1_score(y_test, lr_pred,
average='macro')*100)
print("Weighted Average f1 Score is: ", f1_score(y_test, lr_pred,
average='weighted')*100)
print("f1 Score on Non Weighted score is: ", f1_score(y_test, lr_pred,
average=None)*100)

```

OUTPUT

```

Micro Average f1 Score is:  77.27272727272727
Macro Average f1 Score is:  74.21916104653944
Weighted Average f1 Score is:  76.52373933045479
f1 Score on Non Weighted score is:  [83.09178744  65.34653465]

```

Classification Report of Logistic Regression:

INPUT

```

from sklearn.metrics import classification_report
print('Classification Report of Logistic Regression: \n',
classification_report(y_test, lr_pred, digits=4))

```

Classification Report of Logistic Regression:

	precision	recall	f1-score	support
0	0.7818	0.8866	0.8309	97
1	0.7500	0.5789	0.6535	57
accuracy			0.7727	154
macro avg	0.7659	0.7328	0.7422	154
weighted avg	0.7700	0.7727	0.7652	154

ROC Curve& ROC AUC

```
auc= roc_auc_score(y_test, lr_pred)
print("ROC AUC SCORE of logistic Regression is ", auc)
```


ROC AUC SCORE of logistic Regression is 0.7327726532826913

INPUT

```
from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt

fpr, tpr, thresholds = roc_curve(y_test, lr_pred)
plt.plot(fpr, tpr, color='orange', label="ROC")
plt.plot([0, 1], [0, 1], color='darkblue', linestyle='--', label='ROC curve
(area = %0.2f)' % auc(fpr, tpr))
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("Receiver Operating Characteristics (ROC) Curve of Logistic
Regression")
plt.legend()
plt.grid()
plt.show()
```

OUTPUT



Confusion matrix of "Decision Tree"

Confusion Matrix:

```
from sklearn.metrics import classification_report, confusion_matrix  
  
cm = confusion_matrix(y_test, dt_pred)  
cm
```

OUTPUT


```
array([[83, 14],  
       [15, 42]])
```

INPUT

```
sns.heatmap(confusion_matrix(y_test, dt_pred), annot=True, fmt="d")
```

OUTPUT

<Axes: >



INPUT

```
TN =cm[0, 0]
FP =cm[0,1]
FN = cm[1,0]
TP = cm[1,1]
```

TN, FP, FN, TP

OUTPUT

(83, 14, 15, 42)

INPUT

```
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.metrics import accuracy_score, roc_auc_score, roc_curve
cm = confusion_matrix(y_test, dt_pred)
```

```
print('TN - True Negative {}'.format(cm[0,0]))
print('FP - False Positive {}'.format(cm[0,1]))
print('FN - False Negative {}'.format(cm[1,0]))
print('TP - True Positive {}'.format(cm[1,1]))
print('Accuracy Rate: {}'.format(np.divide(np.sum([cm[0,0], cm[1,1]]),
np.sum(cm))*100))
print('Misclassification Rate: {}'.format(np.divide(np.sum([cm[0,1],
cm[1,0]]), np.sum(cm))*100))
```

OUTPUT

```
TN - True Negative 83
FP - False Positive 14
FN - False Negative 15
TP - True Positive 42
Accuracy Rate: 81.16883116883116
Misclassification Rate: 18.83116883116883
```

INPUT

```
import matplotlib.pyplot as plt
import numpy as np
```

```
plt.clf()
```


```

plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Wistia)
classNames = ['0', '1']
plt.title('Confusion Matrix of Decision Tree')
plt.ylabel('Actual (true) Values')
plt.xlabel('Predicted Values')
tick_marks = np.arange(len(classNames))
plt.xticks(tick_marks, classNames, rotation=45)
plt.yticks(tick_marks, classNames)
s = [['TN', 'FP'], ['FN', 'TP']]
for i in range(2):
    for j in range(2):
        plt.text(j, i, str(s[i][j]) + " = " + str(cm[i][j]))

plt.show()

```

OUTPUT



```
auc = roc_auc_score(y_test, dt_pred)
```

```
print("ROC AUC SCORE of Decision Treeis ", auc)
```


ROC AUC SCORE of Decision Treeis 0.79625610417

INPUT

```
from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt
```

```
fpr, tpr, thresholds = roc_curve(y_test, dt_pred)
plt.plot(fpr, tpr, color='orange', label="ROC")
plt.plot([0, 1], [0, 1], color='darkblue', linestyle='--', label='ROC curve
(area = %0.2f)' % auc(fpr, tpr))
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("Receiver Operating Characteristics (ROC) Curve of Decision Tree")
plt.legend()
plt.grid()
plt.show()
```

OUTPUT



Receiver Operating Characteristic (ROC) curves are a useful tool in an AI-based diabetes prediction system, particularly for assessing the model's ability to discriminate between diabetes and non-diabetes cases. Here's how ROC curves are relevant:

ROC Curve: The ROC curve is a graphical representation that shows the trade-off between the model's true positive rate (recall or sensitivity) and its false positive rate as you vary the classification threshold. The curve is created by plotting these rates for different threshold values.

AUC-ROC: The Area Under the ROC Curve (AUC-ROC) is a common metric used to quantify the overall performance of a model. An AUC-ROC value of 0.5 indicates random guessing, while a value of 1.0 indicates perfect discrimination between the two classes. In the context of a diabetes prediction system, a higher AUC-ROC indicates better discrimination between diabetic and non-diabetic cases.

Threshold Selection: ROC curves help in selecting an appropriate classification threshold based on your system's requirements. A point on the curve that balances sensitivity and specificity can be chosen to optimize model performance.

Model Comparison: You can compare different models or variations of your model by examining their ROC curves and AUC-ROC values. A model with a higher AUC-ROC

generally performs better in discriminating between diabetes and non-diabetes cases.

Imbalanced Datasets: ROC curves are particularly useful when dealing with imbalanced datasets, where one class significantly outnumbers the other. It provides insights into how well the model is handling such imbalances.

➤ ROC curves and AUC-ROC are valuable tools for understanding and evaluating the performance of your diabetes prediction model, especially in terms of its ability to correctly classify patients with diabetes and those without. They provide a visual representation of the model's discriminatory power, which is crucial in healthcare

CONCLUSION:

In conclusion, building an AI-based diabetes prediction system involves several crucial steps, including feature engineering, model selection, and evaluation. Here's a summary of key takeaways for each aspect:

Feature Engineering: Collect and preprocess a diverse dataset with relevant diabetes-related features.

- Handle missing data and outliers appropriately.
- Create new features, if necessary, based on domain knowledge
- Perform feature scaling to ensure features are on a similar scale.

- Conduct feature selection to identify the most relevant predictors.

Model Selection: Choose a variety of machine learning models, including logistic regression, decision trees, random forests, and others.

- Select models suitable for binary classification tasks in a healthcare context.
- Experiment with different algorithms to find the one that performs best on your dataset.
- Consider the interpretability of the model, which is crucial in healthcare applications.

Evaluation: Split the dataset into training and testing sets to assess model performance.

- ☐ Utilize performance metrics like accuracy, precision, recall, F1-score, and AUC-ROC.
- ☐ Be mindful of imbalanced datasets and select appropriate evaluation metrics.
- ☐ Consider the trade-off between precision and recall to balance false positives and false negatives.
- ☐ Use confusion matrices to gain detailed insights into model performance.
- ☐ Explore ROC curves and AUC-ROC to assess the model's discriminatory power.
- ☐ Regularly update and retrain the model as new data becomes available to maintain its predictive accuracy.

In practice, it's essential to iteratively refine your AI-based diabetes prediction system by experimenting with different approaches, models, and evaluation methods. The choice of features, models, and evaluation metrics should align with the specific characteristics of your dataset and the clinical goals of your prediction system. Continuous monitoring and updates are critical to ensure the system remains accurate and reliable over time.